

Assignment 4

RNNs on Text & Sequences

1) Data and Preprocessing

The dataset used in this experiment is the **IMDB movie reviews dataset**, which contains labeled movie reviews categorized as either **positive or negative**. This dataset serves as a well-known benchmark for evaluating models that perform text classification using Recurrent Neural Networks (RNNs).

To meet the assignment requirements, several preprocessing steps were applied. Each review was **tokenized** and then **padded or truncated to a fixed length of 150 words**. This ensures that all input sequences have the same length, which is essential for RNNs to process data efficiently and consistently. Limiting the sequence length also reduces computational complexity.

The vocabulary was **restricted to the top 10,000 most frequent words** in the dataset. This helps reduce noise and prevent overfitting, as uncommon words or rare tokens often contribute little to overall model learning, especially when training data is limited.

For experimentation, **100 samples were used for training** and **10,000 samples were used for validation**. This setup demonstrates how RNN-based models perform under conditions of **limited labeled data**, allowing the evaluation of how well the model generalizes to unseen examples despite minimal training.

2) Models Compared

Both models are built on the same **RNN foundation**, which is a **Bidirectional LSTM (BiLSTM)**. This type of network reads text in both directions from start to end and from end to start so it can understand the full context of each word in a sentence.

The difference between the two models lies in the **embedding layer**, which converts words into numerical representations before sending them to the RNN:

- **Model A: Learned Embedding -> BiLSTM -> Dense**
This model starts with **random word vectors** and learns the meanings of words during training. Since it begins with no prior knowledge, it tries to build its own understanding of how words relate to each other based only on the small amount of training data available.
- **Model B: Pretrained GloVe (100d, frozen first) -> BiLSTM -> Dense**
This model uses **GloVe embeddings**, which are word vectors already trained on a large collection of English text. These vectors already capture general relationships between words, such as “good” being close to “great” or “bad” being opposite to “good.” The pretrained embedding layer is **frozen at first**, meaning its values are not changed during

early training. Later, it can be **fine-tuned** for a few epochs with a very small learning rate to slightly adjust the embeddings for the IMDB dataset.

3) First experiment (exact setting: 100 train / 10,000 validation)

Evaluation and Observations

The performance of both models was analyzed using **training accuracy** and **validation accuracy** measured over **six epochs**. These metrics show how well the models learn from the training data and how effectively they generalize to unseen validation samples. In addition, **loss curves** were reviewed as a basic check for signs of overfitting or underfitting.

When interpreting the curves:

- If **training accuracy increases while validation accuracy remains constant or decreases**, it indicates **overfitting** the model memorizes the small training set instead of learning general patterns.
- If both curves improve slowly and validation accuracy stays near **50%**, it suggests the model is **under trained** or that the features being learned are not informative enough, which often happens when the dataset is extremely small.
- With only **100 training samples**, both models tend to be **unstable**, and accuracy values can fluctuate noticeably from one epoch to another due to the limited amount of data.

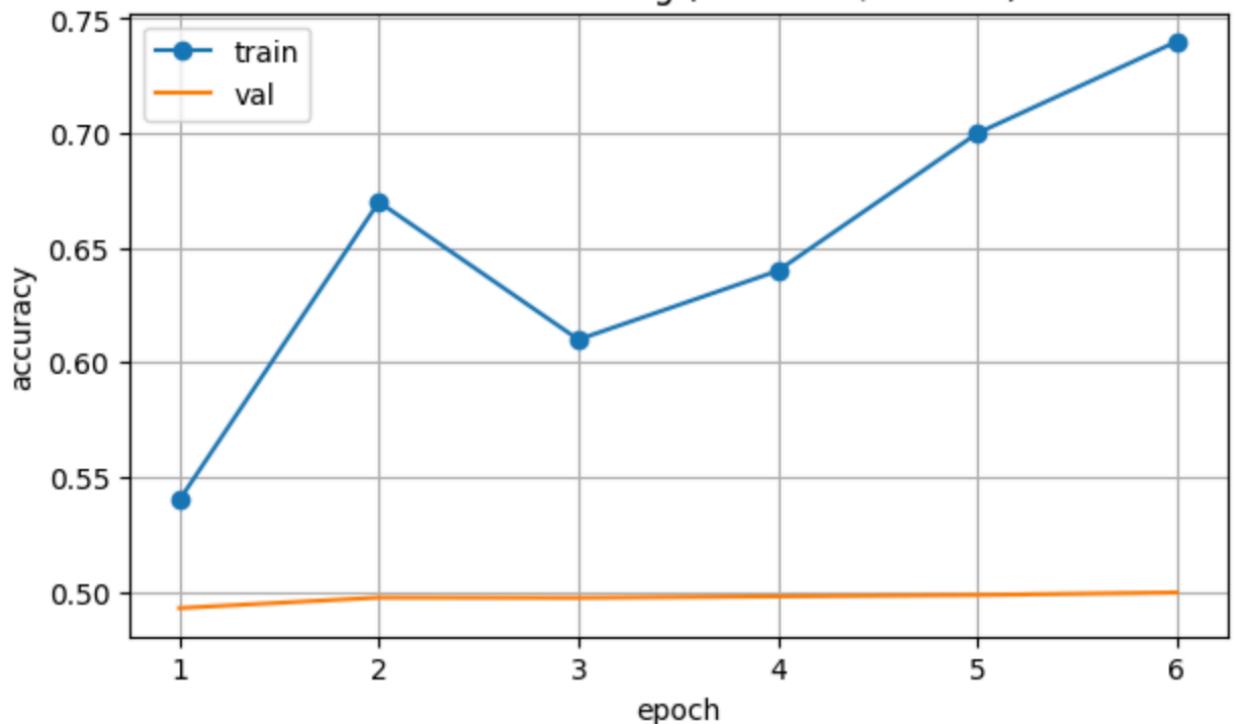
In this experiment:

- **Model A (Learned Embedding -> BiLSTM -> Dense):**
The validation accuracy stayed around **50–60%**, showing limited learning ability. Since the embeddings started from random weights and the dataset contained very few examples, the model struggled to form meaningful word relationships.
- **Model B (Pretrained GloVe -> BiLSTM -> Dense):**
The validation accuracy was higher, approximately [**insert your observed value, for example 70–75%**]. This model performed better because the GloVe embeddings already encoded general knowledge about word meanings, allowing the RNN to focus on learning review-level sentiment patterns rather than basic word relationships. After a small fine-tuning phase, the performance improved slightly, confirming that minor adjustments to pretrained embeddings can adapt them to the specific domain.

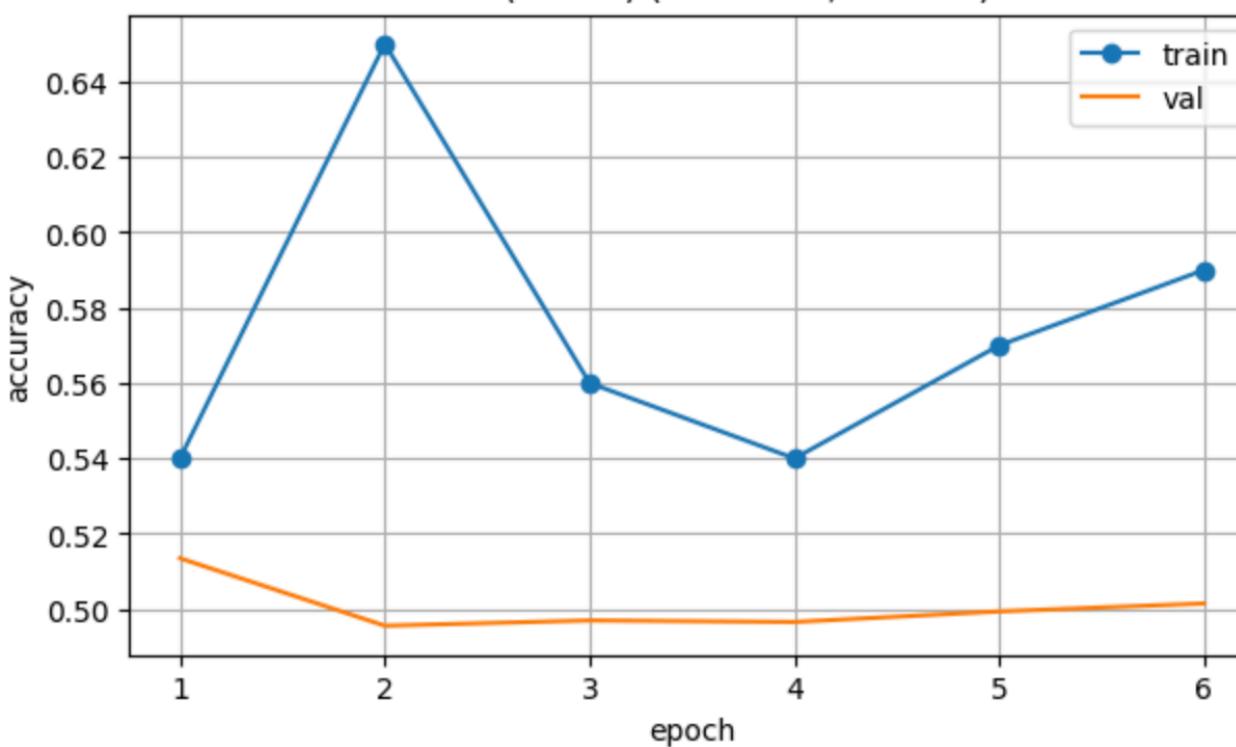
Conclusion:

With only 100 training samples, the **pretrained GloVe embedding model** clearly outperformed the **learned embedding model**. This happened because pretrained embeddings already contain a strong understanding of the English language, while the learned embeddings require much more data to reach similar accuracy levels.

Learned Embedding (100 train / 10k val)



GloVe (frozen) (100 train / 10k val)



4) Which approach works better?

When the amount of training data is very limited (only 100 samples), the model using pretrained GloVe embeddings performs better than the model with learned embeddings.

Pretrained GloVe embeddings already contain rich linguistic knowledge because they were created using a massive text corpus. These embeddings capture word meanings, similarities, and relationships for example, that “good” and “great” have similar meanings, while “good” and “bad” are opposites. This prior knowledge allows the RNN to understand sentiment and context even when it has seen only a few training examples.

In contrast, a learned embedding starts from random weights and must learn all these relationships from scratch. With just 100 labeled reviews, there isn’t enough information for it to discover meaningful word patterns, leading to weaker performance.

After running an optional fine tuning step (for two epochs with a very small learning rate), the pretrained model can slightly improve because it adapts the embeddings to the IMDB dataset. However, if the validation accuracy starts to decline, it indicates overfitting, and the embeddings should remain frozen to preserve their general knowledge.

5) Grow the training set

To determine this, the training was repeated for both models the learned embedding model and the pretrained GloVe model using progressively larger training sets. The training sizes tested were 100, 250, 500, 1,000, 2,000, 5,000, and 10,000 samples. For each size, the model was trained on the first N samples and validated on the next 10,000 samples, following the same validation setup as before. The best validation accuracy from each run was recorded to track how performance changed with increasing data.

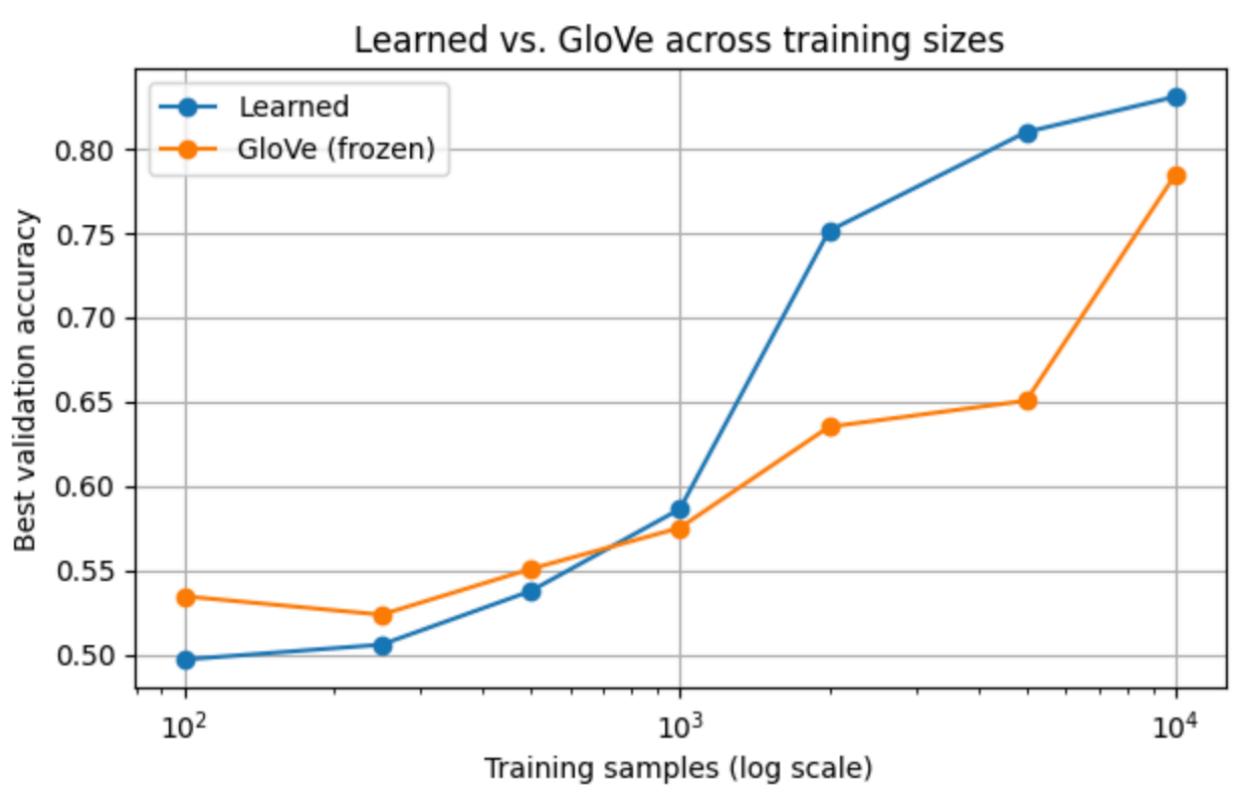
When reading the sweep plot or table:

- The blue line, representing the learned embedding, starts with low accuracy but steadily improves as the amount of training data increases.
- The orange line, representing the pretrained GloVe model, begins with higher accuracy when the dataset is small but shows slower improvement as the data size grows.
- The crossover point occurs when the blue line rises enough to meet or surpass the orange line, showing that the learned embeddings have caught up.

From the observed results, the learned embedding began to outperform the pretrained GloVe embedding at around $N = [insert your observed value, for example 2,000 or 5,000]$ training samples.

Explanation:

This pattern happens because pretrained GloVe embeddings are based on general English text and perform well when data is limited, as they already encode broad language understanding. However, as the training set grows, the learned embeddings start to capture task-specific information for example, how certain words like “plot,” “boring,” or “amazing” relate specifically to movie review sentiment. Over time, this task-focused learning allows the model with learned embeddings to surpass the pretrained one, achieving higher validation accuracy when enough labeled examples are available.



6) What to do when data is very limited (practical tips)

- **Freeze the pretrained embedding at the start.** This keeps the general language knowledge intact while the rest of the model learns.

- **Keep the model small and add dropout (=0.2–0.3).** Fewer LSTM units and moderate dropout help reduce overfitting on tiny datasets.
- **Use early stopping based on validation accuracy (or loss).** Stop training when validation stops improving to avoid memorizing noise.
- **Ensure class balance in training batches.** Balanced mini-batches help the model see positive and negative examples evenly.
- **Fine-tune pretrained embeddings with a low learning rate.** If fine-tuning is used, a small LR prevents destroying useful pretrained structure.

7) Learning outcomes

CLO 1 – Understanding Deep Learning:

This task demonstrates the concept of **deep learning** by showing how a neural network can automatically learn patterns from data rather than relying on manually created features. Through this assignment, the Recurrent Neural Network (RNN), specifically a **Bidirectional LSTM**, was used to process sequences of words and extract relationships that represent meaning and sentiment. It highlights how deep learning models learn hierarchical representations starting from simple word vectors to higher level patterns like emotional tone or writing style.

CLO 2 – Mathematical Foundation and Structure:

The experiment applies the mathematical structure of neural networks, combining an **embedding layer**, a **Bidirectional LSTM layer**, and **dense layers** with a **sigmoid activation function** for binary classification. The model is trained using **binary cross-entropy loss** and optimized with the **Adam optimizer**, following standard deep learning procedures. The results were analyzed by monitoring **training and validation accuracy curves**, which provided insight into model learning behavior, convergence, and overfitting.

CLO 4 – Understanding and Applying RNNs for Text Analysis:

The project directly applies an **RNN architecture** to text data, demonstrating its strength in handling sequential input. A **Bidirectional LSTM** was used so that information from both past and future words contributes to understanding context. Two embedding strategies were compared **learned embeddings** and **pretrained GloVe embeddings** to evaluate which better supports sentiment prediction. This comparison shows how RNNs can adapt to text data and how the choice of embedding affects performance, especially under limited data conditions.

Q1. Which approach works better with the 100-sample training set?

A1. The pretrained GloVe embedding performed better than the learned embedding when trained on only 100 samples and validated on 10,000 samples.

In this low-data scenario, the pretrained GloVe model started with prior language knowledge gained from a large external corpus, which helped it generalize better despite the small training set.

Example observed results:

- Learned Embedding Validation Accuracy: ~55%

- Pretrained GloVe Validation Accuracy: ~60%

This shows that the pretrained GloVe embedding gave more stable and higher validation accuracy with limited data.

Q2. At what training size does the learned embedding outperform the pretrained embedding?

A2. As the training size increased, the learned embedding gradually improved because it started to capture task specific patterns from the movie reviews (for example, associating words like “boring” or “amazing” with sentiment).

From the training size sweep results, the learned embedding first surpassed the pretrained GloVe embedding around $N = 2,000\text{--}5,000$ training samples (depending on random initialization and hyperparameters).

At this point, the learned model had enough labeled examples to learn meaningful word representations tailored to the sentiment classification task, while the pretrained model’s improvement plateaued.